

Pete Gerdzen Project 1: Enron Word Count (Map → Shuffle → Reduce);
Approach: Pure Python (local)

Dataset & Subset

I processed ~517k Enron emails due to local resource limits; results are representative and reproducible via the same selection method.

Results:

- (Top 20) please,137574, would,116308, energy,109133, power,107913, new,107235, said,98068 more,87391, image,84148, time,81827, about,77569, one,77356, gas,76223, which,75917 company,75682, out,73827, thanks,71049, know,69455, information,66297, get,66169 market,66118
- Cumulative tokens in Top-20: 1,725,259
- Vocabulary size: 441,626 unique words

Methods: (aka Describing my approach):

preprocessing (parse_enron.py)

- Prefer text/plain bodies; fallback to text/html with tag stripping.
- Strip quoted replies (drop lines starting with > and sections like “Original Message”).
- Normalize whitespace and remove obvious boilerplate fragments where feasible.
- Deduplicate by SHA-1 of the cleaned body (avoids counting long forward/reply chains repeatedly).
- Write cleaned docs to data/plain_v2/, then point inputs → data/plain_v2.

Word Count Pipeline (main.py) — explicit Map → Shuffle/Combine → Reduce

- MAP: Read each .txt email, lowercase, remove emails/URLs/domains, strip non-ASCII, tokenize on [A-Z/a-z/'], drop tokens ≤2 chars, apply general + Enron-specific stopwords (e.g., *enron*, *ect*, *hou*, months/days). Emit (word, 1) pairs.
- SHUFFLE / COMBINE: For each file, aggregate pairs into per-file term frequencies using plain dict increments.
- REDUCE (two-stage):
 - Merge local counts into a batch accumulator; spill partials to partials/part_XXXX.csv every N files to bound memory.
 - Final global reduce: read all partial CSVs and sum counts into global totals.
- Outputs:

- outputs/top20.csv — the 20 most frequent tokens (header + 20 rows)
- outputs/word_counts.csv — full vocabulary (sorted by frequency)

Reflection: The final counts come from a cleaned, deduplicated subset of **236,244** emails out of **~517,401** raw messages ($\approx 45.66\%$ retained). The **Top-20** terms—*please* (137,574), *would* (116,308), *energy* (109,133), *power* (107,913), *new* (107,235), *said* (98,068), *more* (87,391), *image* (84,148), *time* (81,827), *about* (77,569), *one* (77,356), *gas* (76,223), *which* (75,917), *company* (75,682), *out* (73,827), *thanks* (71,049), *know* (69,455), *information* (66,297), *get* (66,169), *market* (66,118)—together account for **1,725,259** tokens, drawn from a vocabulary of **441,626** unique words. This profile looks like internal corporate email rather than spam: it is dominated by request/coordination language (*please*, *would*, *thanks*, *time*, *get*, *know*) and clearly reflects Enron's domain via sector terms (*energy*, *power*, *gas*, *market*). A few artifacts persist: *image* rises because HTML signatures/logos survive HTML \rightarrow text conversion as placeholders, and thread behavior (forward/reply chains) can repeat prior content even after body-hash deduplication. Generic high-frequency tokens (*one*, *which*, *about*, *out*) remain common in email prose. Overall, the counts point to routine coordination and substantial discussion of energy markets, which is consistent with expectations for the Enron corpus.

Improvement : To push the analysis further, I would first strengthen **thread-aware deduplication** by detecting quoted blocks (From:/Sent:/To: headers and “>” lines) and using Message-ID/References to collapse reply chains so only the newest content segment is counted. Second, I would add **normalization** (lemmatization or careful stemming) so inflected forms merge (e.g., *markets* \rightarrow *market*; *companies* \rightarrow *company*), and slightly expand the **domain stoplist** to exclude residual rendering artifacts (e.g., *image*) and organization-specific boilerplate. Third, I would move beyond unigrams to **phrases and weighting**—compute bigrams/trigrams (e.g., *natural gas*, *power trading*, *credit risk*) and apply **TF-IDF** or sublinear TF so generic words are down-weighted while domain terms and phrases are surfaced. Finally, for scale and clarity, I would keep the current two-stage Reduce but write **time-sliced partials** (e.g., monthly) before the final merge; this preserves reproducibility, enables trend/topic-drift analysis, and makes re-runs faster if only a slice changes.